

EX.NO.: 10

PROBABILITY N-GRAM

DATE: 27.01.2025

To implement bigram modeling on a given corpus (Berkeley restaurant project corpus file) and perform the following tasks:

1. Extract text from the corpus file.
2. Generate all possible bigrams from the text.
3. Apply **Add-One Smoothing** and calculate the probabilities of all bigrams.
4. Query and print the probability of a given random bigram.
5. Implement and apply **Kneser-Ney Smoothing** for the bigram model.
6. Display the probability matrices for both smoothing methods.

PROCEDURE:

1. Read the Text Corpus:
 - a. Use the `read_file()` function to load the corpus from the specified text file and convert it to lowercase for consistency.
2. Tokenize the Text:
 - a. Use NLTK's `word_tokenize()` function to split the text into individual tokens (words).
3. Generate Unigrams and Bigrams:
 - a. Extract unigrams (single tokens) directly from the tokenized text.
 - b. Use the `ngrams()` function from NLTK to generate all bigrams from the tokenized text.
4. Calculate Unigram and Bigram Counts:
 - a. Count the occurrences of each unigram using Counter.
 - b. Similarly, count the occurrences of each bigram using Counter.
5. Add-One Smoothing:
 - a. Implement Add-One smoothing to adjust the probabilities of bigrams, avoiding zero probabilities for unseen bigrams.
 - b. Use the formula: $P(w_2|w_1) = \frac{C(w_1, w_2) + 1}{C(w_1) + |V|}$ where $|V|$ is the vocabulary size.
6. Generate Probability Matrix for Add-One Smoothing:
 - a. Use the `generate_bigram_matrix()` function to construct a matrix where rows and columns represent vocabulary words, and each cell shows the smoothed probability of the bigram.
7. Kneser-Ney Smoothing:
 - a. Implement Kneser-Ney smoothing, which redistributes probability mass to account for lower-order n-gram probabilities.
 - b. Use the formula: $PKN(w_2|w_1) = \frac{C(w_1, w_2) - d}{C(w_1) + d} + \frac{d}{C(w_1) + d} \cdot \frac{C(w_2)}{|V|}$ where d is the discount factor.
8. Generate Probability Matrix for Kneser-Ney Smoothing:
 - a. Use the `generate_bigram_matrix()` function to display the Kneser-Ney smoothed bigram probabilities in a matrix format.
9. Query for a Random Bigram:
 - a. Accept or define a random bigram input (e.g., ('restaurant', 'berkeley')).
 - b. Retrieve and print the probability of the bigram for both Add-One and Kneser-Ney smoothing methods.
10. Output Results:
 - a. Display the bigram probability matrices for both Add-One and Kneser-Ney smoothing.
 - b. Print the queried bigram probabilities with clear labels.

CODE AND OUTPUT

```
import nltk
from nltk.util import ngrams
```

```

from collections import Counter
import pandas as pd

# Function to read text from a file
def read_file(file_path):
    with open(file_path, 'r') as file:
        return file.read().lower()

# Function to calculate bigram probabilities with add-one smoothing
def add_one_smoothing(bigrams, unigram_counts, vocabulary_size):
    bigram_counts = Counter(bigrams)
    smoothed_probabilities = {}

    for bigram in bigram_counts:
        smoothed_probabilities[bigram] = (bigram_counts[bigram] + 1) /
(unigram_counts[bigram[0]] + vocabulary_size)

    return smoothed_probabilities

# Generate a smoothed count matrix
def generate_bigram_count_matrix(tokens, bigrams, vocabulary_size):
    vocab = sorted(set(tokens))
    unigram_counts = Counter(tokens)
    bigram_counts = Counter(bigrams)

    bigram_count_matrix = pd.DataFrame(0, index=vocab, columns=vocab, dtype=int)

    for w1 in vocab:
        for w2 in vocab:
            bigram_count_matrix.loc[w1, w2] = bigram_counts[(w1, w2)] + 1 # Add-One
Smoothing

    return bigram_count_matrix

# Generate a bigram probability matrix
def generate_bigram_probability_matrix(tokens, smoothed_probabilities):
    vocab = sorted(set(tokens))
    bigram_matrix = pd.DataFrame(0, index=vocab, columns=vocab, dtype=float)

    for (w1, w2), prob in smoothed_probabilities.items():
        bigram_matrix.loc[w1, w2] = prob

    return bigram_matrix

# Main function
def main():
    nltk.download('punkt')

    # Read text from file

```

```

file_path = 'corpus.txt' # Replace with your file path
text = read_file(file_path)

# Tokenize the text
tokens = nltk.word_tokenize(text)

# Generate unigrams and bigrams
unigrams = tokens
bigrams = list(ngrams(tokens, 2))

# Calculate unigram counts
unigram_counts = Counter(unigrams)
vocabulary_size = len(unigram_counts)

# Apply Add-One Smoothing
add_one_probs = add_one_smoothing(bigrams, unigram_counts, vocabulary_size)

# Generate bigram count matrix
bigram_count_matrix = generate_bigram_count_matrix(tokens, bigrams,
vocabulary_size)

# Generate bigram probability matrix
bigram_probability_matrix = generate_bigram_probability_matrix(tokens,
add_one_probs)

# Print Bigram Probability Matrix
print("\nBigram Probability Matrix with Add-One Smoothing:")
print(bigram_probability_matrix.round(4))

# Print Bigram Count Matrix
print("Bigram Count Matrix with Add-One Smoothing:")
print(bigram_count_matrix)

# Prompt the user for a bigram input
def get_bigram_probability(smoothed_probabilities, bigram):
    return smoothed_probabilities.get(bigram, "Bigram not found")

# Example usage for querying a bigram
random_bigram = ('spend', 'money') # example input bigram
probability = get_bigram_probability(add_one_probs, random_bigram)
print(f"Probability of the bigram {random_bigram}: {probability}")

if __name__ == "__main__":
    main()

```

Bigram Probability Matrix with Add-One Smoothing:

	chinese	delicious	eat	favorite	food	for	i	\
chinese	0.0000	0.0000	0.0000	0.0000	0.2963	0.0000	0.0000	
delicious	0.1053	0.0000	0.0000	0.0000	0.1053	0.0000	0.0000	
eat	0.2308	0.0000	0.0000	0.0000	0.0769	0.0000	0.0000	
favorite	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
food	0.0000	0.0000	0.1071	0.0000	0.0000	0.1071	0.0714	
for	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
i	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
is	0.0000	0.1429	0.0000	0.0000	0.0000	0.0000	0.0000	
lunch	0.1111	0.0000	0.0741	0.0000	0.0000	0.0000	0.1111	
meal	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
money	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0909	
my	0.0000	0.0000	0.0000	0.1111	0.0000	0.0000	0.0000	
on	0.1000	0.0000	0.0000	0.0000	0.1000	0.0000	0.0000	
spend	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
tasty	0.0000	0.0000	0.0000	0.0000	0.1111	0.0000	0.0000	
to	0.0000	0.0000	0.2800	0.0000	0.0000	0.0000	0.0000	
want	0.0909	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	

	is	lunch	meal	money	my	on	spend	tasty	\
chinese	0.0000	0.1481	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
delicious	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
eat	0.0000	0.1538	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
...									
tasty	1	1	1	1	1	1	1		
to	1	1	1	3	1	1	1		
want	1	1	1	1	1	5	1		

Probability of the bigram ('spend', 'money'): 0.2727272727272727

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings](#)...

[nltk_data] Downloading package punkt to

[nltk_data] C:\Users\Hema\AppData\Roaming\nltk_data...

[nltk_data] Package punkt is already up-to-date!