



```

    return pairs, list(vocabulary)

# Generate positive and negative pairs
pairs, vocabulary = generate_skip_gram_pairs(corpus)

positive_pairs = pairs
negative_pairs = [(random.choice(vocabulary), random.choice(vocabulary)) for _ in
range(len(positive_pairs))]

# Combine positive and negative pairs
all_pairs = positive_pairs + negative_pairs
labels = [1] * len(positive_pairs) + [0] * len(negative_pairs)

# Split into training and test sets
train_pairs, test_pairs, y_train, y_test = train_test_split(all_pairs, labels,
test_size=0.2, random_state=42)

# Convert text pairs into feature vectors
vectorizer = CountVectorizer(analyzer=lambda x: x)
X_train = vectorizer.fit_transform(train_pairs)
X_test = vectorizer.transform(test_pairs)

# Train the logistic regression classifier
clf = LogisticRegression()
clf.fit(X_train, y_train)

# Evaluate the classifier
y_pred = clf.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))

# The classifier can now predict if two words are likely to occur nearby

def predict_proximity(word1, word2):
    pair = [(word1.lower(), word2.lower())]
    pair_vectorized = vectorizer.transform(pair)
    probability = clf.predict_proba(pair_vectorized)[0][1]
    return probability

# Example usage
probability = predict_proximity("stock", "market")
print(f"Probability of 'stock' and 'market' appearing nearby: {probability:.2f}")

```

Accuracy: 0.9129750999788167

	precision	recall	f1-score	support
0	0.89	0.95	0.92	77894
1	0.94	0.88	0.91	77889
accuracy			0.91	155783
macro avg	0.91	0.91	0.91	155783
weighted avg	0.91	0.91	0.91	155783

Probability of 'stock' and 'market' appearing nearby: 0.81

### INFERENCE

The Logistic Regression classifier successfully learns the proximity relationship between words based on Skip-Gram pairs. A higher accuracy indicates that the model effectively captures the semantic relationship between words in the corpus.